



Extended Steering Anticipation — at Scale

Anticipation only pays AT SCALE: steering on a noisy single-well forecast backfires; a fleet-learned dip prior is what makes it win.


One well: naive anticipation (no prior) drifts out of zone; prior-anchored ties reactive




Fleet of 200 wells
(mean % footage in-zone)



Look-ahead projection adds noise — best at ~0 ft horizon



The real lever: the fleet/sheaf dip prior — 71%→98% in-zone



Finding: anticipation pays ONLY through a reliable, fleet-learned dip prior (the sheaf/structural model, this repo): in-zone climbs 71%→98% as the prior strengthens, and the fleet mean beats reactive (90%→95%).

Anticipating off single-well data is WORSE than reacting; extending the look-ahead horizon does NOT help (best ~0 ft).